

TEU00311

What is the Internet doing to me? (witidtm)

Stephen Farrell
stephen.farrell@cs.tcd.ie

<https://github.com/sftcd/witidtm>
<https://down.dsg.cs.tcd.ie/witidtm>


URLs accessed 20201019

Computers, Operating Systems and some example risks

Overall landscape

- Desktops, laptops, tablets, phones, “smart” speakers, home routers, raspberry-Pis, “fitness” devices...are all computers running some operating system
- Differently popular, open, reliable, invasive and general purpose or not
- Your choices should not only be driven by money and ease-of-use, though those factors tend to dominate

Operating Systems Market Shares



Graph: last 10 years

Android
Windoze
iOS
OS X

...
Desktop Linuxes
e.g Ubuntu

<https://gs.statcounter.com/os-market-share#quarterly-200901-202004>

Linux (laptop/desktop) now <1%

Machine/OS issues to consider (1)

- Price, ease of use, updates/support, backup
 - Insurance? Meh;-)
- Security: probability of attack, probability of success, cost (impact) of successful attack
 - Probability of attack relates to market share
- Privacy: hiding in crowds vs. standing out by being more privacy-aware
 - Tension vs. Probability of attack?
- Having backups is your #1 protection – if you do nothing else set that up!
 - But only if you sometimes test restore!
- Automated updates should be considered mandatory, or else you'll eventually suffer
 - Consider how well update system works, some OSes take ages to update
 - Some devices (still!) get no updates – avoid!

Machine/OS issues to consider (2)

- Openness – to what extent can someone find out what's going on under the hood?
- Even if you can't/don't do that, whether or not others can makes a difference
- It's a big world and many of the tech issues you face will have been overcome by someone else sometime
 - And likely being nerdy, they won't have been able to resist telling the world how they fixed stuff:-)
- But: “open” tends to be correlated with “hard to use” or “for nerds”

Machine/OS issues to consider (3)

- To be fair, there are also some benefits in a more complete but closed ecosystem, like Apple's
 - https://manuals.info.apple.com/MANUALS/1000/MA1902/en_US/apple-platform-security-guide.pdf
 - Note1: the above document is written for someone like me, it may be a tough read;-) But it is very thorough.
 - Note2: the link above may be subject to bit-rot – the 2019 version no longer worked in 2020 but did sort-of land me on the right bit of Apple web page
- Interesting to compare/contrast today's "major" OSes – Windows, Android, Apple with what the FOSS movement would like
 - FOSS: https://en.wikipedia.org/wiki/Free_and_open-source_software

End of Life

- Devices do not last forever
 - Ultimately the battery has a limited lifetime in terms of charge cycles
 - When (when, not if!) a device fails or is lost/stolen, what will you lose?
 - See “backup” point earlier
 - When (when, not if!) a vendor turns off support, or a service (e.g. OS updates, some photo-sharing feature), what will happen to your stuff?
 - When you’re done with the hardware, what do you do?
 - “Two-thirds of used disc drives on Craigslist and eBay contain sensitive data” is a 2016 article, no reason to expect that’ll change v. soon
- <https://www.itpro.co.uk/security/26814/two-thirds-of-used-disc-drives-on-craigslist-and-ebay-contain-sensitive-data>

Device-based tracking/surveillance

- Anything with “location services” can be unexpectedly dodgy
 - **Image metadata inside the image file can include location**
 - “Fitness” trackers
- Cell-tower history within mobile operators
- List of Wi-Fi networks to which you’ve attached sometime
- **Bluetooth Tracking**
- “Smart” speakers
- “Security” cameras
- Telemetry from applications or OS
- Software update!
 - But didn’t we want that? Yes. But can’t it help tracking? Yes. Hmm...

Blackboard Assignment Heads-Up

- Research and report on:
 - 1) the ways in which a real device/OS you possess or use exposes you to potential device-based tracking, as you normally use the device/OS
 - 2) how the browser(s) you use accumulate cookies and other state and how you could reduce or avoid that
- Report: ~4-5 pages, more if needed but try keep it short, use URLs/references/appendices for tech detail or screen shots
- Content:
 - For (1): say how that tracking could work, describe your exposure, say what you think about that and what you will or will not do about it - justify that last part, esp. If you conclude you'll do nothing
 - For (2): we'll chat after we learn a bit more about Advertising
- Deadline: Before reading week (ish, meh, deadlines;-)
- Marks: 40%, higher marks for the unexpected and for thorough research!
- We'll talk about two tracking examples, you can't use those for the assignment
- More on this later, feel free to ask questions.... esp. In Monday "office-hour"

Image Metadata: EXIF

<https://photographylife.com/what-is-exif-data>

<https://helpdeskgeek.com/how-to/how-to-remove-exif-data-from-your-photos/>




```
exif DSCN0010.jpg
EXIF tags in 'DSCN0010.jpg' ('Intel' byte order):
-----
Tag                |Value
-----
Image Description   |
Manufacturer        |NIKON
Model               |COOLPIX P6000
Orientation         |Top-left
X-Resolution        |300
Y-Resolution        |300
Resolution Unit     |Inch
Software            |Nikon Transfer 1.1 W
Date and Time       |2008:11:01 21:15:07
YCbCr Positioning   |Centered
Compression         |JPEG compression
X-Resolution        |72
Y-Resolution        |72
Resolution Unit     |Inch
Exposure Time       |1/75 sec.
F-Number            |f/5.9
Exposure Program    |Normal program
ISO Speed Ratings    |64
Exif Version        |Exif Version 2.2
Date and Time (Orig)|2008:10:22 16:28:39
Date and Time (Digit)|2008:10:22 16:28:39
Components Configure|Y Cb Cr -
Exposure Bias       |0.00 EV
Maximum Aperture Val|2.90 EV (f/2.7)
Metering Mode       |Pattern
Light Source        |Unknown
Flash               |Flash did not fire, compulsory flash mode
```

```
Focal Length        |24.0 mm
Maker Note          |3298 bytes undefined data
User Comment        |
FlashPixVersion     |FlashPix Version 1.0
Color Space         |sRGB
Pixel X Dimension    |640
Pixel Y Dimension    |480
File Source         |DSC
Scene Type          |Directly photographed
Custom Rendered     |Normal process
Exposure Mode       |Auto exposure
White Balance       |Auto white balance
Digital Zoom Ratio   |0.00
Focal Length in 35mm|112
Scene Capture Type   |Standard
Gain Control        |Normal
Contrast            |Normal
Saturation          |Normal
Sharpness           |Normal
Subject Distance Ran|Unknown
North or South Latit|N
Latitude            |43, 28,
2.81400000
East or West Longitu|E
Longitude           |11, 53,
6.45599999
Altitude Reference   |Sea level
GPS Time (Atomic Cl)|14:27:07.24
GPS Satellites       |06
GPS Image Direction  |
Geodetic Survey Data|WGS-84
GPS Date            |2008:10:23
Interoperability Ind|R98
Interoperability Ver|0100
-----
EXIF data contains a thumbnail (6702 bytes).
```

<https://raw.githubusercontent.com/ianare/exif-samples/master/jpg/gps/DSCN0010.jpg>

52100 Arezzo, Province of Arezzo, Italy



- Took about 5 minutes to find this in Google street view
- Most of that was finding a way to map degree, minutes, seconds to fractional Lat,Long
- All **automatable**, could easily produce location history from a set of images
- How could such a “leak” be damaging to you or to someone else in your images?

Assignment Part-1 Outline (1)

- If you were to use location-in-EXIF for the assignment (which you can't:-) then your assignment might look like:
 - Intro para: summarise all below like a paper abstract
 - Para describing the general vulnerability: device tracking as a risk to device-owner
 - Para describing EXIF, GPS, how **your** camera uses that (or used to, if that's the case)
 - Describe potential for tracking: uploaded images giving a location history
 - What'd you do about it?
 - Default to location-services off, scrub images before upload (say how), check earlier uploaded images (and do what with 'em?), get a new device/camera
 - Or, if your response is “do-nothing,” then justify that, e.g. I don't upload images but get lost a lot so need location-services:-) Or maybe argue that image locations aren't a big deal except in unrealistic movie-plot scenarios
 - References – if reference is to peer-reviewed paper, good; if reference is URL then include “accessed-at date”

Assignment Part-1 Outline (2)

- Try make it something you'd not be embarrassed to show your buddies in a year's time
- It's ok if the tracking you describe is low-risk, i.e., you do not have to find some mega-concerning risk to write about
- It's ok if you reference something where you don't fully grok all the details but if so, say so and we can chat about that in class if you want
- **PLEASE** submit a PDF
- It's ok to work together on this one, but do say with whom you've worked but the conclusions must be your own
- Effort:
 - 4-5 pages is fine, as is more
 - Maybe spend an hour researching (20 mins searching, 40 reading) and another hour writing for Part-1
- Possible starting point: Google scholar (want a bit of a demo?)
 - https://scholar.google.com/scholar?hl=en&as_sdt=0%2C5&q=device+tracking+image+exif+location+data&btnG

MAC Addresses

- Device-tracking often (ab)uses **long-term hard-coded identifiers** such as MAC addresses (or IMEI/IMSI in mobile n/w)
- MAC address: layer 2 address (mostly) hardcoded to radio or other network chip
 - Same form of address used in WiFi and most other network protocols at layer 2, e.g. Bluetooth
 - Roughly how two devices on the same local area network (LAN) identify one another
- Looks like “6C:9C:ED:87:27:60” (48 bits) - 1st half is manufacturer ID (Cisco), 2nd half device-ID (a WiFi router in TCD SCSS)
 - You can look up the manuf ID in the registry: <https://www.adminsub.net/mac-address-finder/84:C7> gives a list that includes Sony
- MAC address is often fixed for the lifetime of the device; There is now a 64-bit version, not sure how widely used yet
 - You can probably see these in the “about device” tab or similar
- MAC address randomisation is a good idea and starting to be deployed
 - Often, the MAC address only really needs to be stable for a session, so can be randomised
 - But – if you paid for the hotel WiFi that might be based on your MAC address, or an enterprise network might use MAC addresses to decide which machines are allowed on the local network, or the machine may be a switch/router/server where changing MAC address would break stuff or be inefficient
 - So you can't always randomise, and doing so well needs higher-layer controls

Bluetooth

- Early BT devices regularly broadcast their MAC addresses
- Later BT specs try to fix this, allowing support for randomised MAC addresses
- All good so? Nope!
 - Even if MAC addresses and payload identifiers change randomly, doing so out of sync enables tracking, sometimes indefinitely (read paper, and see next slide)
 - Becker, J., Li, D., & Starobinski, D. (2019). “Tracking Anonymized Bluetooth Devices,” Proceedings on Privacy Enhancing Technologies, 2019(3), 50-65. doi: <https://doi.org/10.2478/popets-2019-0036>
- Generalising: identifier correlation over time, devices and networks, is a very hard problem to mitigate – similar issues with DHCP leases, IP and email addresses, account names and web artefacts (cookies etc.)




Fig. 9. An experiment illustrating the carry-over effect on Windows 10 devices. Asynchronous value changes allow updating the device identity whenever it changes.




Fig. 10. This timeline shows address-carryover across 5 random addresses on iOS. The first three hops occur via the handoff identifying token, the last one occurs via the nearby token. Different colors denote different values.

Malicious Software (Malware)

- Malware on your devices could:
 - Exfiltrate your data (keyboard sniffers, ID theft, “wallet” attacks, mobile app “SDKs”)
 - Work for someone else (cryptominers, clickfraud/adware)
 - Be part of a botnet (DDoS)
 - Encrypt your files (ransomware)
 - (Sometimes) survive factory-reset (though less commonly)
- Technical types: viruses, worms, trojans, ...
 - We’re less concerned here with how malware works or details of how to detect/prevent malware doing bad things
 - More interesting is: How can you avoid or recover?
- Note that anti-malware tools (e.g. anti-virus) can look similar to malware - they both want access the “innards” of the OS/network-interactions – that can make such tools attractive as a target

How malware may arrive...

- Pre-installed (accidentally or deliberately, bloatware/crapware)
- Supply chain attacks (e.g., buy a popular npm library, infect that, or generate a new library with a typosquatted name)
- Appear to be an application or browser extension that someone may want (e.g., phone torch app)
- Mail attachment (pdf, zip, tar, ...) in spam or spear-phish – be careful of how email addresses and file names are displayed in your mail user agent!
- Drive-by exploitation via web-site/browser vulnerability
- USB stick (watering hole attack)
- USB charger (how much do you trust Dublin Bus? :-)

Actors in the “AppStore” model

- The interested parties include the OS vendors, device manufacturers, mobile network operators, app developers and the services those apps use/enable
 - Your interests likely come last in that list, despite what most of the others may say, and even if they really mean what they say!
 - Their interests won't always be aligned either (e.g. mobile network operator vs. over-the-top service provider)
 - And there are non-stupid bad actors in the game, esp. in the app developer and services-those-use/enable cohorts
- So what do you think are their interests?
 - OS makers: Google (android), Apple (IOS), Canonical (Ubuntu)
 - Device manufacturers: Samsung, Apple, Huawei, Sony, ...
 - Mobile network operators: China Mobile, ... Vodafone, ... Telefonica, ... Eir, Three
 - App developers: Google, Apple, device makers plus many others I've never heard of
 - Service providers: Google, Facebook, Twitter, ... Netflix, Sky, ...

“AppStore” Permissions model

- Both iOS and android “app” install processes use a “permissions” model, windoze, chromebooks and MAC OS are heading that way more
 - As is Ubuntu, but only sort-of, it being a much more open OS
- Android used (before 2015) present “permissions needed” only at install time, now both it and iOS pop up permission requests when the app first tries to make use of something “sensitive”
 - But are those permission classes really meaningful? And do they protect you? TBH: I find it hard to tell
- Research in this space also seems less concentrated on how well all this protects users and more on how “well accepted” it is by users, e.g.:
 - Reinfelder, Lena, et al. "An Inquiry into Perception and Usage of Smartphone Permission Models." International Conference on Trust and Privacy in Digital Business. Springer, Cham, 2018.
 - https://www.researchgate.net/profile/Lena_Reinfelder/publication/326630389_An_Inquiry_into_Perception_and_Usage_of_Smartphone_Permission_Models_15th_International_Conference_TrustBus_2018_Regensburg_Germany_September_5-6_2018_Proceedings/links/5c0e58564585157ac1b73e03/An-Inquiry-into-Perception-and-Usage-of-Smartphone-Permission-Models-15th-International-Conference-TrustBus-2018-Regensburg-Germany-September-5-6-2018-Proceedings.pdf
 - Rajivan, Prashanth, and Jean Camp. "Influence of privacy attitude and privacy cue framing on android app choices." Twelfth Symposium on Usable Privacy and Security ({SOUPS} 2016). 2016.
 - https://www.usenix.org/system/files/conference/soups2016/wpi16_paper-rajivan.pdf
 - Could be I didn't spend enough time looking though (I did look around in 2019 but didn't repeat in 2020 so far:-)

What do I use?

- Main devices:
 - Home router: OpenWRT (Turris Omnia)
 - Laptop: Ubuntu (Dell XPS)
 - Phone: Sailfish (Sony Xperia)
 - A cardboard box: with a pile of old phones and laptop hard-drives:-)
- Various others @ home and for work, mostly Ubuntu servers on hosted virtual machines (VMs), some real (“bare-metal”) servers and some old laptops
- Other people on my home network use android, iphone, mac, windoze
- Main living room laptop also runs Ubuntu and is entirely usable for non-nerds

If you wanna try Ubuntu...

- Ideally try with old laptop
 - Drivers for newer models can take a while to catch up with proprietary versions, esp for WiFi
- You can try before installing to hard drive by running from a “live” USB stick
- You can make system dual-boot if you wanna try some more but still keep your earlier OS
- I wish we had some face-to-face sessions where you can try stuff like that out... but we don't. If the situation improves in the new year, and a bunch of you wanted to, I'd be happy to organise a bit of a hackathon like that.
 - If so: mail me **after** you've gotten your module results

What do you use?

What would you like to use?

What do you dislike that you use?