

# TEU00311

What is the Internet doing to me?  
(witidtm)

Stephen Farrell  
[stephen.farrell@cs.tcd.ie](mailto:stephen.farrell@cs.tcd.ie)

<https://github.com/sftcd/witidtm>  
<https://down.dsg.cs.tcd.ie/witidtm>

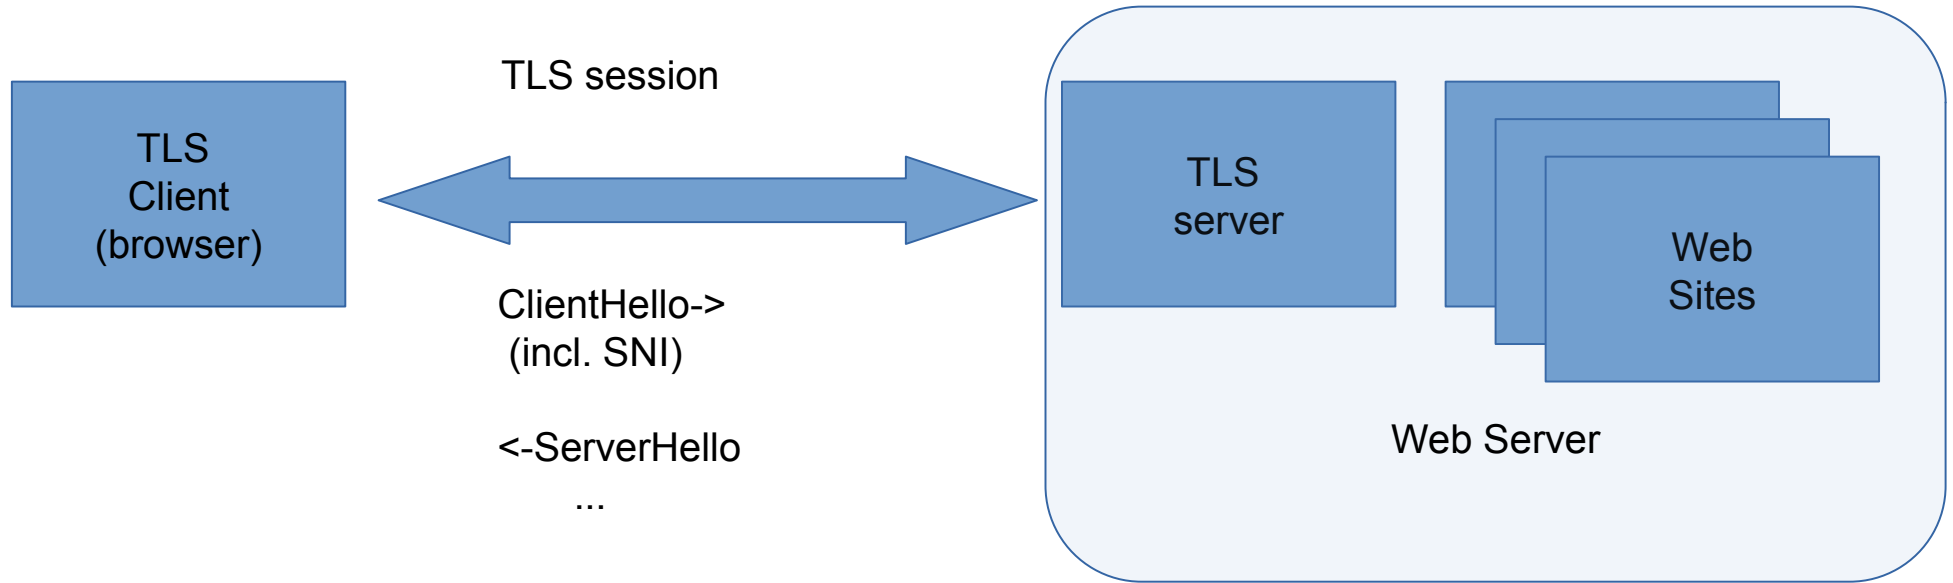
# Overview

- As well as knowing about today's Internet it might be useful to know a bit about what's coming and how some things evolve
- So I'll describe some work I've been doing for the last couple of years, how that might go and how it might ultimately affect you
- That's a thing called "Encrypted ClientHello"

# TLS and SNI

- Transport Layer Security (TLS, RFC8446) is the security protocol that secures the web and many other applications – HTTP running over TLS is what makes HTTPS
- One web server instance (e.g. an apache install using VirtualHost) can, and very frequently does, serve multiple web sites
- Each of those may (and is v. likely to) use different TLS server key pairs/certificates, which are the things that allow a TLS client (like your browser) to authenticate the web site, i.e. to know that you're really connecting to someone who (is related to someone who) controls "tcd.ie"

# TLS for multiple web sites

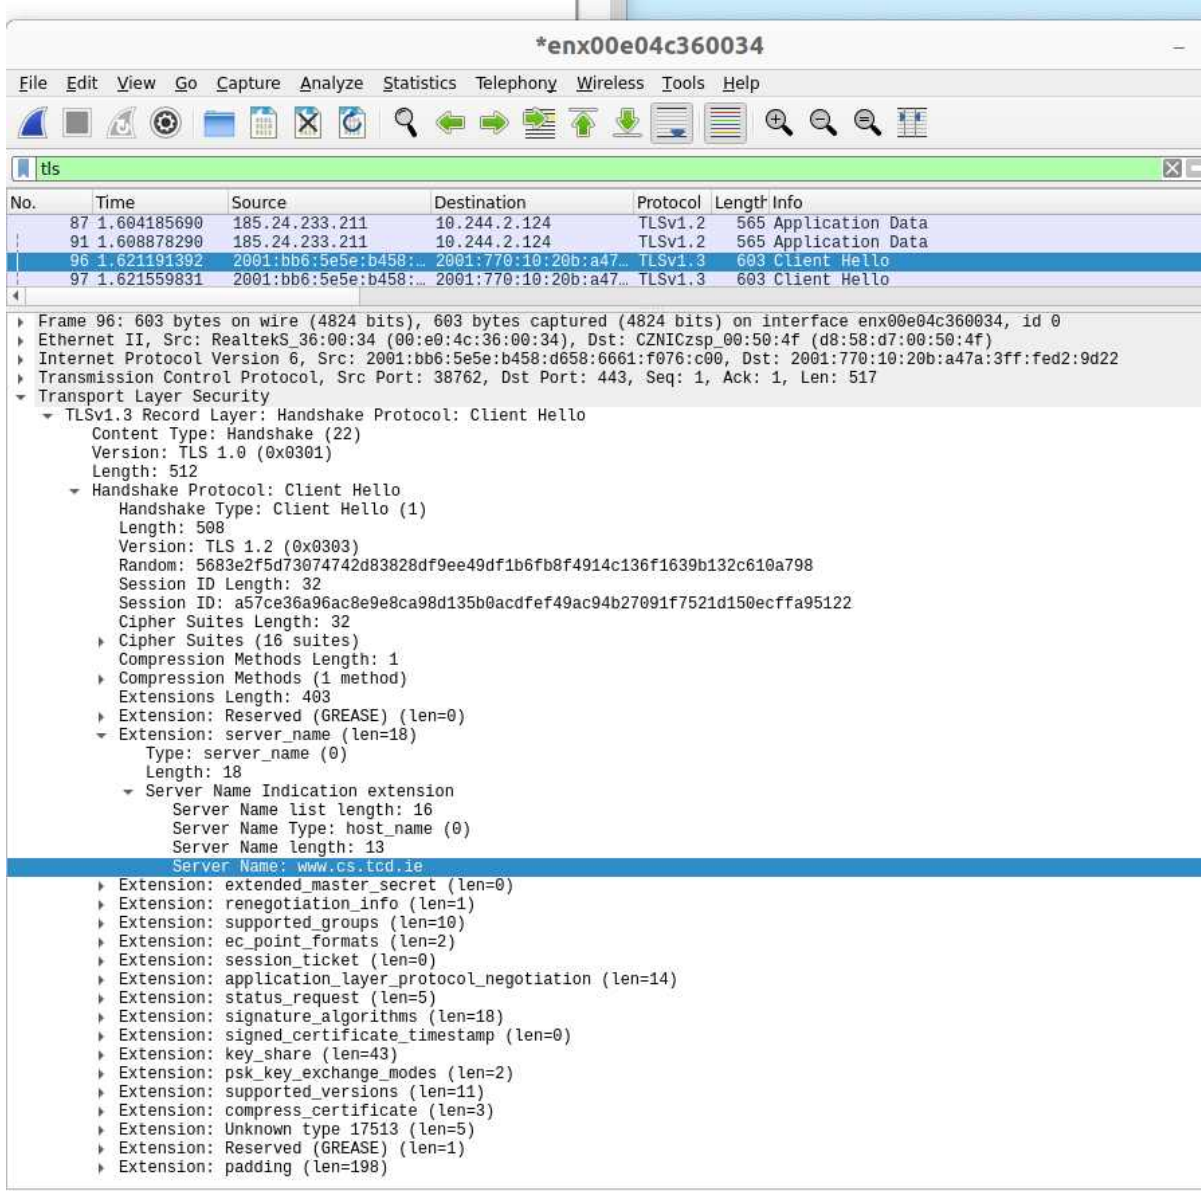


# TLS and SNI

- One of the first things an HTTPS server needs to do is pick a TLS server key/certificate to use for the TLS session
- The client needs to verify that it's talking to the correct server via the TLS server certificate, which (for the web) contains the domain name of the web site
- Result: the first TLS message the client sends (the ClientHello) needs to specify the web site (DNS name) for which the TLS session is being established
- That's done using the Server Name Indication (SNI) extension to the ClientHello (RFC6066)

# The SNI “Leak”

- Remember Wireshark?
- It knows how to decode this kind of thing and we can see the cleartext SNI value on the right
- That ClientHello message was sent from a browser when I accessed <https://www.cs.tcd.ie>



# SNI as a leak

- Since the SNI value is sent in the first message, neither party has a key with which to encrypt, so SNI is sent in clear in the first TLS handshake message (the ClientHello)
- When accessing <https://bank.example.com/getBalance> the “getBalance” part will be encrypted (later, when the full HTTP request is sent) but the DNS name (“bank.example.com”) is sent in clear in the SNI extension
- That’s a noticeable leak, especially as the SNI is visible to everyone on the path (my ISP, the site’s ISP, every intermediate router, hosters, governments)
- SNI has also been used for censorship

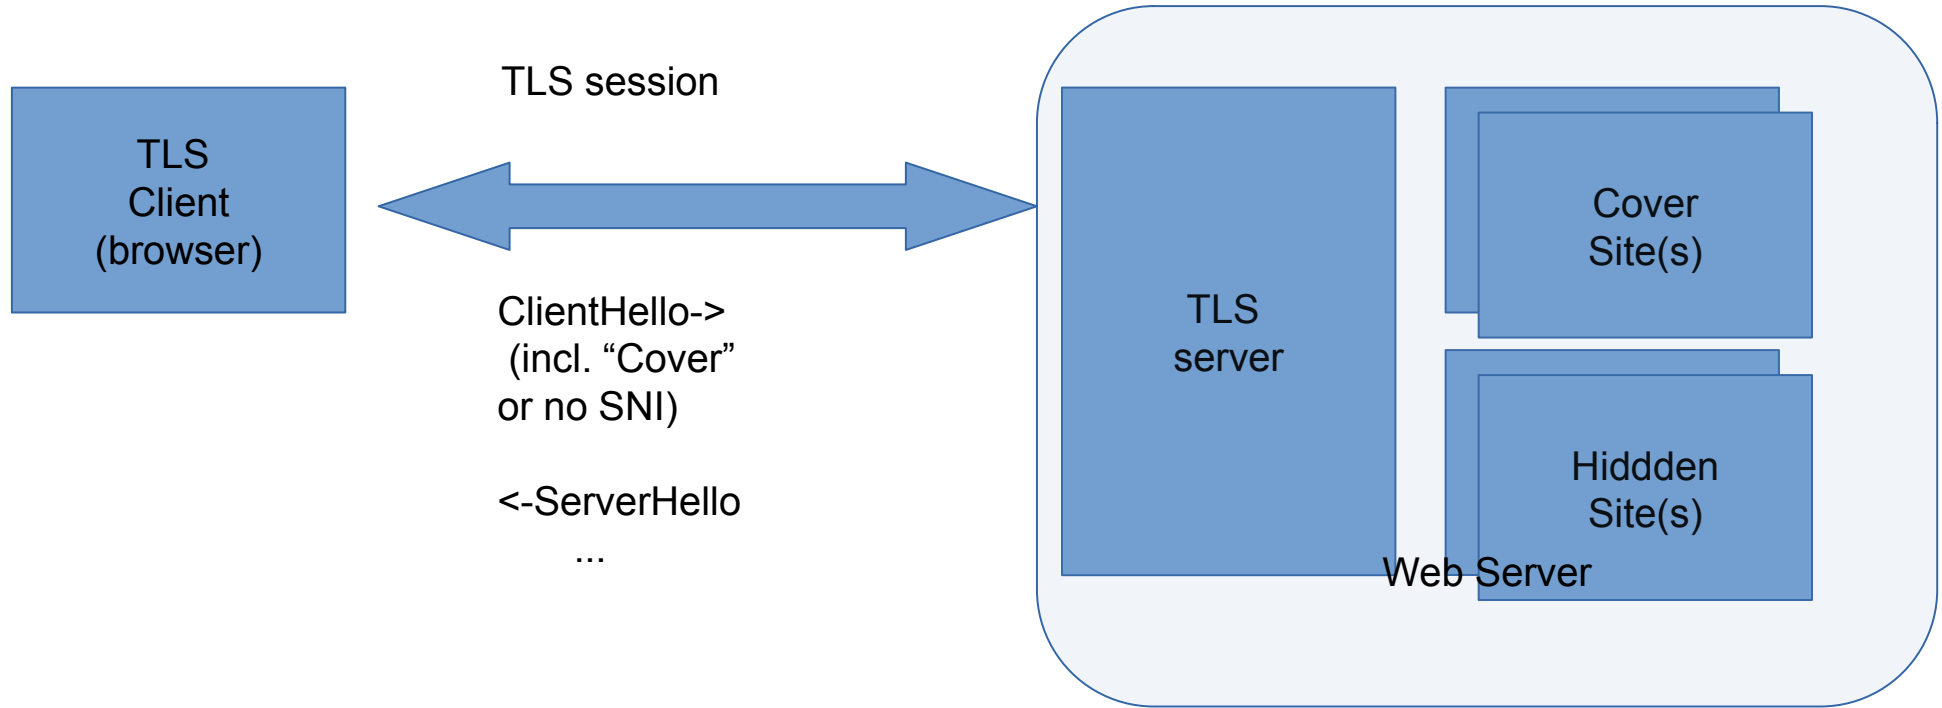
<https://www.bleepingcomputer.com/news/security/south-korea-is-censoring-the-internet-by-snooping-on-sni-traffic/>

# SNI as a leak

- Domain “fronting” (where the SNI has the hoster’s name but the HTTP request has the “real” DNS name) was brittle and got turned off by service providers
- Previously, we weren’t motivated to try address this quite tricky problem because...
  - TLS server certificate was sent in the clear prior to TLS1.3
  - DNS name sent in clear using DNS protocol, but now we do have DNS privacy mechanisms (DoT/DoH)



# What we'd like, and can do now ("co-located" variant)



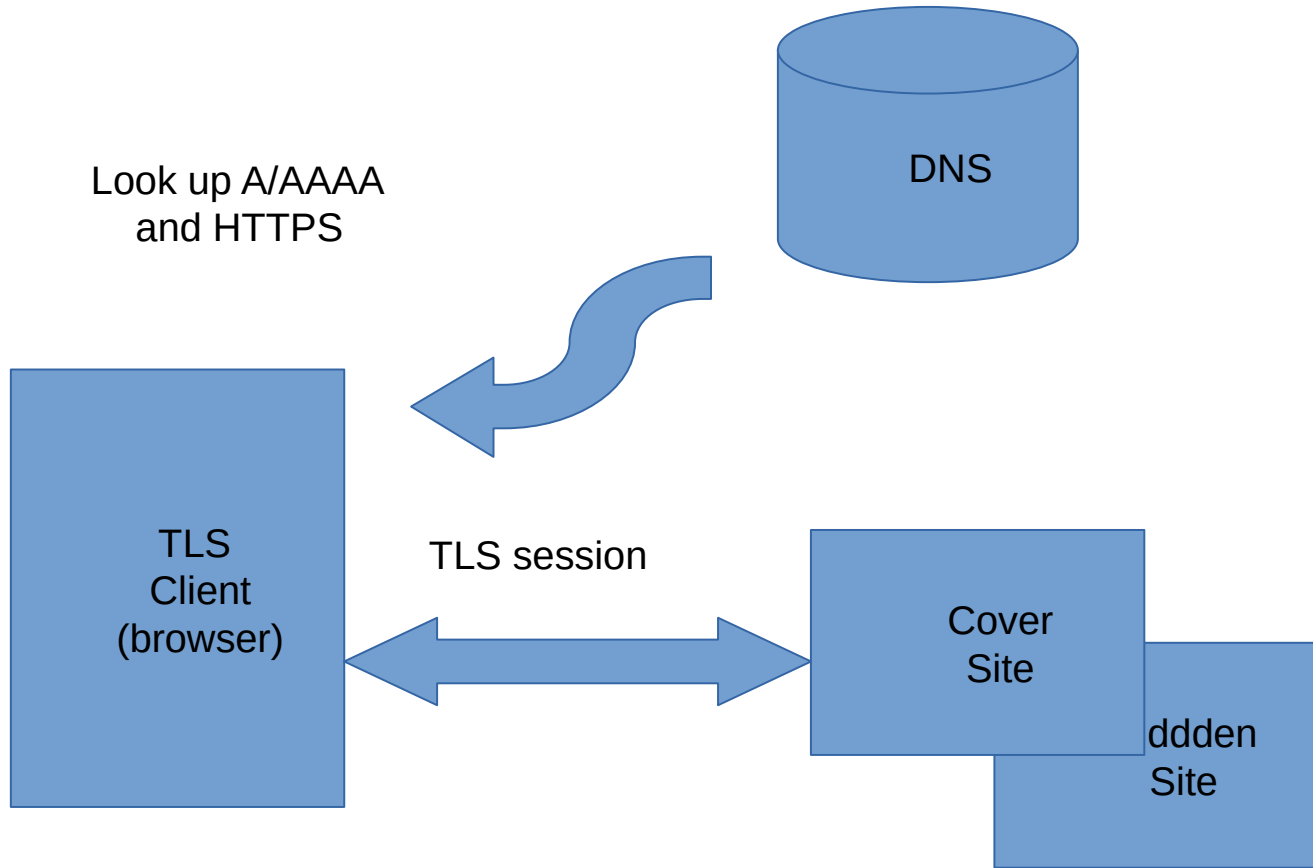
# Encrypted ClientHello (ECH)

- Solution being developed in the IETF TLS WG:  
<https://tools.ietf.org/html/draft-ietf-tls-esni>
- Current draft is version -13, been in development since 2018
- Latest version seems complete, now at the stage of testing
- Multiple implementations exist, including mine
  - <https://github.com/sftcd/openssl/> is my “fork” of OpenSSL - a very widely used TLS/encryption library that can be used with e.g. apache, nginx, lighttpd (web server implementations)

# How does ECH work?

- Needs ability to create/consume new DNS resource records
- Needs TLS1.3 (earlier versions send server cert in clear)
- DNS privacy (DoT/DoH) not strictly needed but if you don't use that maybe there's less point in using ECH (Browsers will likely couple the two)
- Web site publishes a new public key/value in the DNS ("SVCB" or "HTTPS") with some additional keys for ECH
- ECH-aware client (e.g. browser) can check if DNS record exists and has ECH keys
- All going well, use those ECH keys to derive a new shared-secret and send the "real" ClientHello message encrypted inside an "outer" ClientHello message
- The fact that ECH is being used is still visible

# ECH Picture



# GREASEing ECH

- The fact that ECH is being used is still visible
- That may be countered via “GREASEing” - having browsers that are not using ECH sometimes send a ClientHello that looks like it does use ECH
- GREASEing is an anti-ossification TLS implementation trick – clients and servers include garbage values for optional things in order to decrease the probability that middleboxes fixate on currently deployed protocol options - RFC8701

# Interesting thing #1

- People who hate DNS privacy, hate ECH even more; for them, this is browsers and major web sites taking away information from which they've benefited
  - They used to monitor DNS queries but now DoT/DoH make that much harder
  - Switched to monitoring SNI from TLS ClientHello messages, and now ECH is killing that off
- They may have been using that information for what you consider good or bad (they considered it “good”):
  - Censorship, net-nanny, corporate policy enforcement, whitelisting TLS sessions to not MITM
- That DNS or SNI information could also be used for profiling or sold to advertisers, but I'm not aware of reliable information that that has happened

# Interesting thing #2

- The people proposing/backing ECH are (I believe) doing so to try improve privacy (e.g. me, mozilla, aclu)
- ECH may however further increase centralisation since hiding in larger crowds is more effective
- It could be that the likes of cloudflare (test server deployed), apple (implemented client and server earlier) and google (implemented client so far) are the main beneficiaries of ECH – how should we consider that result?

# Interesting thing #3

- ECH should make the kind of measurement Doug described harder, maybe a lot harder
- How do we try get accountability for software and systems without leaving open gaping holes to be exploited by attackers or nation-states?
  - Answer: we don't know (or at least I don't)



# Interesting thing #4

- My guess is that whether or not, and how well, Google chrome implement ECH GREASEing will be the key determinant of whether or not ECH gets long term deployment
- As ECH sticks out, censors like the GFW can just block it
- That's easy if it's not widely used and would kill deployment
- If almost all ClientHello messages appear to contain ECH then such blocking gets harder
  - Maybe not for the GFW but perhaps for most other censors

# Last interesting thing

- The processes for defining and implementing ECH are open to anyone with the skills and interest to get involved
  - Consumes a lot of time, but getting the skills required isn't really hard, all of you are certainly smart enough
- So you could (if you choose) end up involved in key decisions about how the Internet affects you that can reverberate up to fairly highly political levels
  - (Don't get a big head though:-)
  - That's part of the "permissionless innovation" I mentioned at the start
- It's also possible to be as involved with fewer technical skills, e.g. in civil society organisations or campaigning groups

# Last slide for the module

- I hope you enjoyed the module, learned stuff and aren't too depressed about it all
- I'm really sorry that the technical community of which I've been part has produced this crap surveillance capitalism stuff
- But there are alternatives and practices you can adopt with not too much thought or work – do that! You are not at all near helpless!
- I also hope this annoys you sufficiently that you think about how to do something more about the bits you consider unacceptable

# Thanks!

- Feel free to mail me anytime (next year, whenever) if there's relevant stuff with which I can help